

Learning from User-generated Data

Summer Term 2022

Learning from Explicit User Feedback: Memory-based Collaborative Filtering



Markus Schedl

markus.schedl@jku.at

<http://www.cp.jku.at>



Recommender systems are omnipresent

For products:



For fashion:



For jokes:

Jester 5.0

Jokes for *your* sense of humor

For travel:




For movies/series:

NETFLIX

For books:



For music:



For users and
user-generated
content:



Retrieving, Browsing, Recommending

Definitions and Implications:

Information **retrieval** (IR) is the area of study concerned with **searching** for documents, for information within documents, and for metadata about documents [...] *[Wikipedia]*

Implication: user has a **specific information need**, typically represented as a query; system tries to find the most relevant items wrt. query

Retrieving, Browsing, Recommending

Definitions and Implications:

Browsing: interactive task in which the user is more interested in **exploring** the document collection than in retrieving documents which satisfy a specific information need *[Baeza-Yates and Ribeiro-Neto, 1999]*

Implication: user has an **undirected information need** (hard to express as a query); system should provide a good interface for discovering items

Retrieving, Browsing, Recommending

Definitions and Implications:


Recommender: form of a specific type of information filtering system technique that attempts to **recommend** information items [...] or social elements [...] that are likely to be of interest **to the user**. *[Wikipedia]*

Implication: the system searches for potentially relevant items based on user's actions or preferences that form a **user profile** (can, e.g., be observed during retrieval or browsing)

Retrieving, Browsing, Recommending

Required Involvement of the User:

Retrieval: user has a specific information need, typically represented as a query; system tries to find the most relevant items wrt. query



user needs to be active

Browsing: user has an undirected information need; system should provide a good interface for discovering items

Recommending: the system searches for potentially relevant items based on user's actions or preferences (e.g., observed during retrieval or browsing)



**user can be
(almost) passive**


Characteristics of Recommender Systems: RS...

- provide a **personalized** view on the available collection (personalized filtering): every user sees a different list depending on her taste
- need to know something about the user and maintain some sort of **user model** or **user profile** - generated either **implicitly** (e.g., by observing behavior or preferences or keeping a list of purchased items) or **explicitly** (e.g., by asking to rate items)
- **find** those items that are most likely of interest to the user and (possibly) **rank** them according to prospective interest (utility)
- to this end often incorporate **additional knowledge** such as the user's characteristics (e.g., demographic filtering) or context (e.g., context-aware recommendation)
- **need to scale** well: typically deal with very large databases (millions of items) and communities (up to millions of users)

Main Flavors of Recommendation Systems


Collaborative filtering:

Recommend to target user items that other *similar users* liked in the past



Content-based filtering:

Recommend to target user *content similar* to what he or she liked in the past



Context-aware RS:

Recommend to target user items that he, she, or other users liked in a given *context or situation*



Hybrid RS: Any *combination* of the above

Multimedia Mining

User Modeling

Why People Use Recommender Systems

Herlocker et al. have identified 10 reasons why people use recommender systems (user tasks):

- 1. Annotation in Context.** Predictions of potential items related to current context are displayed (historical background: recommending discussion postings, links while Web browsing)
- 2. Find Good Items.** “Core recommendation task”; Interfaces suggests top-ranked list and shows predicted rating values (optional)
- 3. Find All Good Items.** In contrast to 2., in this case, systems must not overlook a single relevant item, i.e., keep false negative rate low (scenario: law case precedents, patents)
- 4. Recommend Sequence.** Whole sequence should be pleasant or meaningful, cf. automatic playlist generation in personalized radio streams or reading recommendations for scientific literature to learn about a field

Why People Use Recommender Systems

5. **Just Browsing.** Browsing (commercial systems) can be pleasant even without the goal of purchasing; pure entertainment
6. **Find Credible Recommender.** “Playing around” with the system to see whether recommendations match taste; checking system for *bias* by using multiple profiles
7. **Improve Profile.** Rating items to improve user profile to improve recommendations
8. **Express Self.** Contributing ratings and reviews is satisfactory for some people (sometimes even more important than getting recommendations)
9. **Help Others.** To let the community benefit from own contributions
10. **Influence Others.** Manipulate ratings to make others view or buy (or avoid) specific items, e.g., *hacking* and *vandalism*

Collaborative Filtering (CF)

- Most widely used approach for recommendation
- Applicable to many domains (no explicit domain knowledge necessary)
- Main underlying assumption: users that had similar taste in the past, will have similar taste in the future
- Collect user feedback given as explicit (or implicit) ratings of items
- Use ratings of all or some users collectively to determine whether to recommend a new item to a user (e.g., by predicting the user's rating for an item)

Gathering Explicit User Data: Rating Methods

Different types (scales) of rating data:

- “Continuous” slider and/or text box



- Likert scale: discrete steps (e.g., 5 or 10 stars, possibly in half steps)



- Binary ratings (thumbs up/down)



- Unary ratings (“Like”)



(For investigation of impact of different rating scales, cf. [Sparling and Sen, 2011])

User-Item-Matrix / Rating Matrix

- Representation of ratings in a user-item matrix:

$U = \{u_1, \dots, u_n\}$... set of users

$P = \{p_1, \dots, p_m\}$... set of items

R ... rating matrix ($n \times m$), $r_{i,j}$ represents rating of user i for item j

	<i>Item 1</i>	<i>Item 2</i>	<i>Item 3</i>	<i>Item 4</i>	<i>Item 5</i>
<i>User 1</i>	3		2	3	3
<i>User 2</i>	4	3	4	3	
<i>User 3</i>	3	2	1	5	4
<i>User 4</i>		5	4	3	1

- Task: predict missing rating (for item 5) for another active user

	<i>Item 1</i>	<i>Item 2</i>	<i>Item 3</i>	<i>Item 4</i>	<i>Item 5</i>
<i>User a</i>	5		3	4	?

CF: Model-based vs. Memory-based

Memory-based CF:

- item ratings made by users are stored
- no model is built, but predictions are made on the fly
- nearest neighbor approach (find most similar users or items)
- simple and well-established method
- results easy to interpret → transparency
- very memory-demanding

Model-based CF:

- user-item matrix is factorized
- ratings are explained by “latent factors” in a low-dimensional space
- users and items can be represented in the same space
- once model is created, rating prediction is very fast
- not so simple, but well-established method
- latent factors may not be easy to interpret (just describe variance in the data)

CF: User-based vs. Item-based

User-based CF:

- find similar users based on rated items (vectors over items)
- predict rating as weighted combination of most similar users' ratings

Item-based CF:

- find similar items based on user ratings (vectors over users)
 - predict rating as weighted combination of most similar items' ratings
-
- In real-world applications, item-based methods can scale better.
 - Rating biases should be accounted for.

User-Based CF Recommendation

- Idea:
identify **similar users**, use their ratings to predict missing rating
- Algorithm outline:
 1. Calculate similarity of active/target user to all users that have rated the item to predict
 2. Select k users that have highest similarity (*k-neighborhood*)
 3. Compute prediction for item from a weighted combination of the item's ratings of users in neighborhood
(weights correspond to similarity)

User-Based CF Recommendation

1. Calculate similarity (=weight) of active user to all users that have rated the item to predict
-

- Commonly used for user similarity: **Pearson's correlation**

$$\text{sim}(a, u) = \frac{\sum_{p \in P'} (r_{a,p} - \bar{r}_a)(r_{u,p} - \bar{r}_u)}{\sqrt{\sum_{p \in P'} (r_{a,p} - \bar{r}_a)^2} \sqrt{\sum_{p \in P'} (r_{u,p} - \bar{r}_u)^2}}$$

where P' is the set of items rated by both users and \bar{r}_u is the mean rating of user u :

$$\bar{r}_u = \frac{1}{|P'|} \sum_{p \in P'} r_{u,p}$$

- Ranges from -1 to $+1$, requires variance in user ratings (else undefined), accounts for *users' rating biases* (general high or low ratings) by subtracting mean rating

User-Based CF Recommendation

1. Calculate similarity (=weight) of active user to all users that have rated the item to predict

- Pearson's correlation has shown to work best for this purpose
- Alternatives are *Spearman rank correlation* (e.g., to deal with different numeric ranges of ratings), *Kendall's τ correlation*, *mean squared differences*, *entropy*, etc.

2. Select k users that have highest similarity (*neighborhood*)

- Predefine k , sort according to similarity scores, and select k highest (should not need any further explanation...)

User-Based CF Recommendation

3. Compute prediction for item from a weighted combination of the item's ratings of users in neighborhood
-

- Predict rating r' as weighted average of deviations from neighbors' mean rating on target item

$$r'_{a,p} = \bar{r}_a + \frac{\sum_{u \in K} \text{sim}(a,u) * (r_{u,p} - \bar{r}_u)}{\sum_{u \in K} \text{sim}(a,u)}$$

where K is the set of the k nearest neighbors and \bar{r}_a the mean rating of the active user a (this time calculated over all of a 's ratings)

- Starts from a 's rating bias and adds deviations based on similarity
- After predicting all missing values of a , the items with highest prediction will be recommended to a

User-Based CF Recommendation – Example

Back to our example...

- User 2 hasn't rated item 5...

1. Calculate correlations

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	3		2	3	3
User 2	4	3	4	3	
User 3	3	2	1	4	4
User 4		5	4	3	1
User a	5		3	4	?

$$\begin{aligned}
 \text{sim}(a, u_1) &= \frac{(5-4)(3-2.67) + (3-4)(2-2.67) + (4-4)(3-2.67)}{\sqrt{(5-4)^2 + (3-4)^2 + (4-4)^2} \sqrt{(3-2.67)^2 + (2-2.67)^2 + (3-2.67)^2}} \\
 &= \frac{0.33 + 0.67 + 0}{\sqrt{2} \sqrt{0.11 + 0.44 + 0.11}} = \frac{1}{1.15} = 0.87
 \end{aligned}$$

$$\begin{aligned}
 \text{sim}(a, u_3) &= \frac{(5-4)(3-2.67) + (3-4)(1-2.67) + (4-4)(4-2.67)}{\sqrt{(5-4)^2 + (3-4)^2 + (4-4)^2} \sqrt{(3-2.67)^2 + (1-2.67)^2 + (4-2.67)^2}} \\
 &= \frac{0.33 + 1.67 + 0}{\sqrt{2} \sqrt{4.66}} = \frac{2}{3.05} = 0.65
 \end{aligned}$$

We will ignore all users that are negatively (or un-) correlated!

$$\text{sim}(a, u_4) = \frac{(3-3.5)(4-3.5) + (4-3.5)(3-3.5)}{\sqrt{(3-3.5)^2 + (4-3.5)^2} \sqrt{(4-3.5)^2 + (3-3.5)^2}} = \frac{-0.25 - 0.25}{0.5} = -1$$

User-Based CF Recommendation – Example

2. Sort and select neighbors
(for the setting $k=2$):

i.e., $K = \{u_1, u_3\}$

	Item 1	Item 2	Item 3	Item 4	Item 5
User 1	3		2	3	3
User 2	4	3	4	3	
User 3	3	2	1	4	4
User 4		5	4	3	1
User a	5		3	4	?

3. Calculate prediction for item 5 for user a

$$r'_{a,i_5} = 4 + \frac{[0.87 * (3 - 2.75)] + [0.65 * (4 - 2.8)]}{0.87 + 0.65} = 4 + \frac{0.9975}{1.52} = 4.66$$

Thus, we predict a rating of **4.66** (or 4.5 or 5, depending on the scale)

Is this a good prediction?

What would be the predicted rating for item 2?

And which of the two would you recommend to user a ?

→ optional homework, exercise :)

Advanced User-Based CF Recommendation

- Users' agreement on items that are “controversial” should have more impact than agreement on items everybody likes
 - *inverse user frequency*: items liked by all get less influence
 - *variance weighting factor*: items with high variance in ratings get more influence
- Neighborhood selection should also consider how many items have been co-rated (small overlap might lead to high correlation by chance)
 - *significance weighting*: devalue weights if less than 50 co-rated items
 - *default voting*: assume default values for missing ratings to calculate correlations on union of rated items rather than on intersection
- General problem: proper selection of k (or a similarity threshold)
 - *coverage* vs. *noise* ($k < 10$ should be avoided, typically $20 < k < 50$)

Item-Based CF Recommendation

- Idea: identify **similar items** based on all their user ratings; use ratings active user gave to those similar items to predict missing rating
 - Algorithm very similar to user-based CF recommendation (instead of the rows in the rating matrix, we now look at the columns)
1. Calculate similarity of items by comparing rating vectors (over users)
 - instead of correlation, calculate (*adjusted*) *cosine similarity*
 2. Select k items that have highest similarity (*k-neighborhood*)
 3. Compute prediction for item from a weighted combination of the ratings of items in neighborhood (weights correspond to similarity)


Item Similarity

Cosine similarity:

$$\text{sim}(i, j) = \cos \theta = \frac{\vec{i} \cdot \vec{j}}{|\vec{i}| * |\vec{j}|} = \frac{\sum_{u \in U} r_{u,i} * r_{u,j}}{\sqrt{\sum_{u \in U} r_{u,i}^2} \sqrt{\sum_{u \in U} r_{u,j}^2}}$$

where $i, j \in P$, \vec{i}, \vec{j} their corresponding rating vectors,
 • the dot product of vectors, and $|\vec{i}|$ the Euclidian length of \vec{i}
 θ is the angle between the vectors, U users that rated both i and j

- Range: 0 to 1 (1 indicates max. sim.) when all rating values are positive
- Can't deal with missing values
- Does account for different item popularities (due to normalization)
- Does not account for user rating bias!
 → adjusted cosine sim.



Item Similarity 2 and Rating Prediction

Adjusted cosine similarity:

$$sim(i, j) = \frac{\sum_{u \in U} (r_{u,i} - \bar{r}_u)(r_{u,j} - \bar{r}_u)}{\sqrt{\sum_{u \in U} (r_{u,i} - \bar{r}_u)^2} \sqrt{\sum_{u \in U} (r_{u,j} - \bar{r}_u)^2}}$$

where U is the set of users that rated both i and j .

- cf. Pearson correlation; range also from -1 to $+1$
- Subtracts user u 's average from ratings given to items i and j
(Not item average! This is not a Pearson correlation on R^T !)

Rating prediction: weighted sum of user's ratings for similar items

$$r'_{a,p} = \frac{\sum_{n \in N} sim(n, p) * r_{a,n}}{\sum_{n \in N} sim(n, p)}$$

where N is the set of the k most similar items

Item-Based Recommendation – Example

Again our example...

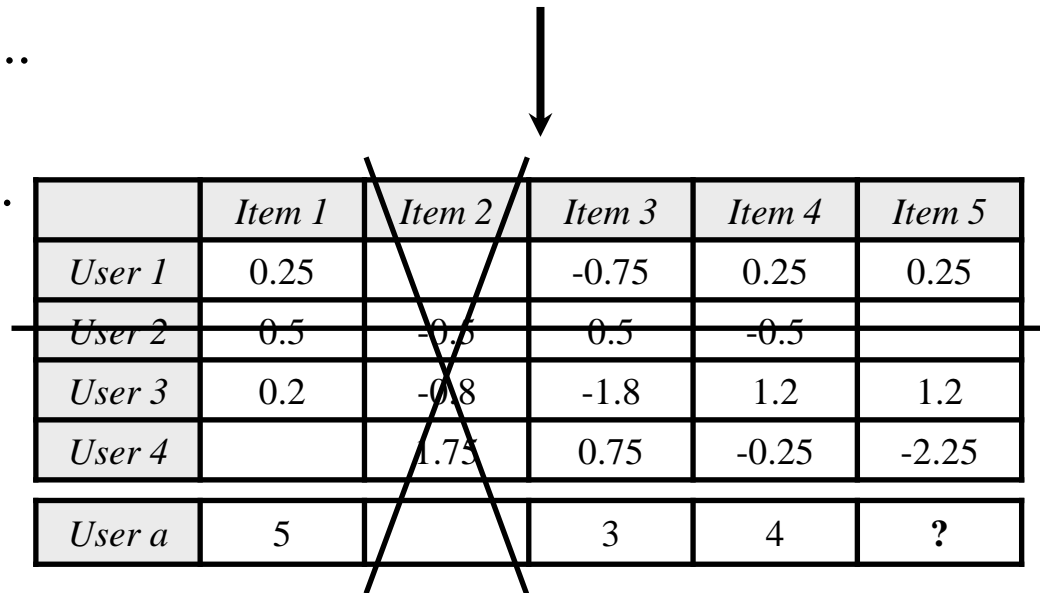
	<i>Item 1</i>	<i>Item 2</i>	<i>Item 3</i>	<i>Item 4</i>	<i>Item 5</i>
<i>User 1</i>	3		2	3	3
<i>User 2</i>	4	3	4	3	
<i>User 3</i>	3	2	1	4	4
<i>User 4</i>		5	4	3	1
<i>User a</i>	5		3	4	?

1. Subtract mean user rating
(saves some time later)

2. User 2 won't help for item 5...

And user *a* hasn't rated item 2...

3. Calculate adjusted cosine
similarities
(i.e., cosine similarity on
mean-adjusted matrix)



	<i>Item 1</i>	<i>Item 2</i>	<i>Item 3</i>	<i>Item 4</i>	<i>Item 5</i>
<i>User 1</i>	0.25		-0.75	0.25	0.25
<i>User 2</i>	0.5	-0.5	0.5	-0.5	
<i>User 3</i>	0.2	-0.8	-1.8	1.2	1.2
<i>User 4</i>		1.75	0.75	-0.25	-2.25
<i>User a</i>	5		3	4	?

Item-Based Recommendation – Example

(Compacted matrix)

3. Adjusted cosine similarities

	Item 1	Item 3	Item 4	Item 5
User 1	0.25	-0.75	0.25	0.25
User 3	0.2	-1.8	1.2	1.2
User 4		0.75	-0.25	-2.25
User a	5	3	4	?

$$\text{sim}(i_5, i_1) = \frac{(0.25 * 0.25) + (1.2 * 0.2)}{\sqrt{0.25^2 + 1.2^2} \sqrt{0.25^2 + 0.2^2}} = 0.77$$

$$\text{sim}(i_5, i_3) = \frac{(0.25 * -0.75) + (1.2 * -1.8) + (-2.25 * 0.75)}{\sqrt{0.25^2 + 1.2^2 + 2.25^2} \sqrt{0.75^2 + 1.8^2 + 0.75^2}} = -0.68$$

$$\text{sim}(i_5, i_4) = \frac{(0.25 * 0.25) + (1.2 * 1.2) + (-2.25 * -0.25)}{\sqrt{0.25^2 + 1.2^2 + 2.25^2} \sqrt{0.25^2 + 1.2^2 + 0.25^2}} = 0.82$$

4. Select $k=2$ most similar (i_1, i_4) and calculate rating prediction

$$r'_{a, i_5} = \frac{0.77 * 5 + 0.82 * 4}{0.77 + 0.82} = 4.48$$

Item-Based CF Recommendation

- Better suited for large-scale recommenders than user-based CF
- Preprocessing can be performed offline, i.e., all *item-to-item similarities* can be calculated in advance (need update after some time)
(Could be done for user-to-user similarities too, but user profiles, esp. of new users, change very dynamically)
- More realistic: users rate only small number of items ($\ll m$)
To predict item i , find most similar (item-item similarity matrix lookup), and weight own ratings over these (commonly only few) items
(In contrast, user-based CF needs to consider all users who rated the target item)
- For item-based CF, at runtime, recommendation in real-time possible

Popular Real-World Data Sets

Name	Domain	Users	Items	Ratings	Sparsity
BX	<i>Books</i>	278,858	271,379	1,149,780	0.9999
EachMovie	<i>Movies</i>	72,916	1,628	2,811,983	0.9763
Jester	<i>Jokes</i>	73,421	101	4.1M	0.4471
MovieLens 10M	<i>Movies</i>	71,567	10,681	10M	0.9869
Netflix	<i>Movies</i>	480K	18K	100M	0.9884
Yahoo! Music	<i>Music</i>	1M	625K	262M	0.9995

Sparsity: Fraction/percentage of missing ratings in R

Summary: Memory-based CF

- Memory-based CF recommenders store all ratings made by users and calculate a prediction on the fly
- Very memory-demanding; not the most effective (caching strategies required), item-based CF typically more effective
- Make use of nearest neighbor approaches (“**lazy learning**”)
- Memory-based CF recommenders are easily interpretable, well-established, and researched
- Simple method that is prone to several negative effects (data sparsity, “cold-start”, curse of dimensionality, etc.)
- Solutions to counteract these effect exist: graph-based transitivity, mutual proximity, etc.

References: Textbooks, Tutorials

[Ricci et al., 2015] *Recommender Systems Handbook*, 2nd ed., Springer.

[Aggarwal, 2016] *Recommender Systems: The Textbook*, Springer.

[Jannach et al., 2010] *Recommender Systems - An Introduction*, Cambridge Press, 2010.

[Baeza-Yates and Ribeiro-Neto, 1999] *Modern Information Retrieval*, Addison Wesley.

[Jannach and Friedrich, 2011] *Tutorial: Recommender Systems*, International Joint Conference on Artificial Intelligence,
<http://ijcai-11.iiia.csic.es/files/proceedings/Tutorial%20IJCAI%202011%20Gesamt.pdf>

References: Original Research

- [Herlocker et al., 2004] *Evaluating collaborative filtering recommender systems*, ACM Transactions on Information Systems, 22(1).
- [Jannach and Friedrich, 2011] *Tutorial: Recommender Systems*, International Joint Conference on Artificial Intelligence,
<http://ijcai-11.iia.csic.es/files/proceedings/Tutorial%20IJCAI%202011%20Gesamt.pdf>
- [Lemire and Maclachlan, 2005] *Slope One Predictors for Online Rating-Based Collaborative Filtering*, In SIAM Data Mining.
- [Linden et al., 2003] *Amazon.com recommendations: item-to-item collaborative filtering*. In IEEE Internet Computing, 7(1).
- [Melville and Sindhvani, 2010]. *Recommender Systems*, In Encyclopedia of Machine Learning, Springer.
- [Resnick et al., 1994] *GroupLens: An Open Architecture for Collaborative Filtering of Netnews*, In Proceedings of the CSCW.
- [Sparling and Sen, 2011] *Rating: how difficult is it?*, In Proceedings of the 5th ACM Conference on Recommender Systems.